Representation Learning

Natural Language Processing

University of Maryland

Update Examples

Imports

```python
import numpy as np
import torch
```

Simple Regression

```python
x = torch.tensor(0.)
w = torch.tensor(2., requires_grad=True)
b = torch.tensor(30., requires_grad=True)
# If you forget "requires_grad", expect this error:
# RuntimeError: element 0 of tensors does not require

def forward(x):
    return w * x + b
```

Try it out!

- Try to predict from input of 20

Inputs and Outputs

```
inputs = torch.tensor([[-40],
                        [-10],
                        [15],
                        [30],
                        [50]])

targets = torch.tensor([[-40],
                        [14],
                        [59],
                        [86],
                        [122]])
```

Inputs and Outputs

```
inputs = torch.tensor([[-40],
                       [-10],
                       [15],
                       [30],
                       [50]])

targets = torch.tensor([[-40],
                        [14],
                        [59],
                        [86],
                        [122]])
```

What are we predicting? What are predictions on **inputs**?

Current Predictions

```
tensor([[-50.],
        [ 10.],
        [ 60.],
        [ 90.],
        [130.]], grad_fn=<AddBackward0>)
```

Current Predictions

```
tensor([[-50.],
        [ 10.],
        [ 60.],
        [ 90.],
        [130.]], grad_fn=<AddBackward0>)
```

What's the MSE loss of these predictions?

Loss Function

```python
def mse(t1, t2):
    diff = t1 - t2
    return torch.mean(diff**2)

loss = mse(preds, targets)

>>> print(loss)
tensor(39.4000, grad_fn=<MeanBackward0>)
```

Loss Function

```python
def mse(t1, t2):
    diff = t1 - t2
    return torch.mean(diff**2)

loss = mse(preds, targets)

>>> print(loss)
tensor(39.4000, grad_fn=<MeanBackward0>)
```

Next: create backpropagation signal to w and b!

```
loss.backward()
print("Backprop signal to w:")
print(w)
print(w.grad)

print("Backprop signal to b:")
print(b)
print(b.grad)
```

Gradients

```
Backprop signal to w:
tensor(2., requires_grad=True)
tensor(390.)
Backprop signal to b:
tensor(30., requires_grad=True)
tensor(-0.4000)
```

Gradients

```
Backprop signal to w:
tensor(2., requires_grad=True)
tensor(390.)
Backprop signal to b:
tensor(30., requires_grad=True)
tensor(-0.4000)
```

What is this saying about *w* and *b*?

- Correct formula is

$$f = \frac{9c}{5} + 32 \tag{1}$$

- *w* should be smaller and *b* should be bigger
- Learning rate and batch size are important, trust Pytorch to do better job!

Learning rate details and multi-objective optimization

- Correct formula is

$$f = \frac{9c}{5} + 32 \tag{1}$$

- *w* should be smaller and *b* should be bigger
- Learning rate and batch size are important, trust Pytorch to do better job!
- Update parameters

Parameter updates

```
with torch.no_grad():
    w -= w.grad * 1e-4
    b -= b.grad * 1e-4

w.grad.zero_()
b.grad.zero_()
```

Parameter updates

```python
with torch.no_grad():
    w -= w.grad * 1e-4
    b -= b.grad * 1e-4

w.grad.zero_()
b.grad.zero_()
```

What are predictions and loss now?

New predictions and loss

```
New predictions:
tensor([[-48.4400],
        [ 10.3900],
        [ 59.4150],
        [ 88.8300],
        [128.0500]], grad_fn=<AddBackward0>)
New loss:
tensor(25.8098, grad_fn=<MeanBackward0>)
```